

Lab 3: Monitoring Process Behavior and Dynamic Priority Scheduling [260 pts]**Due: 10/18/2023 (Wed.), 11:59 PM****1. Objectives**

The objective of this lab is to extend XINU's fixed priority scheduling to dynamic priority scheduling that facilitates fair sharing of CPU cycles among processes. First, we add instrumentation code and new system calls so that CPU usage and response time of app processes is monitored. Second, we compensate I/O-bound processes that voluntarily relinquish CPU before depleting their time slice by improving responsiveness compared to processes that hog a CPU. We do so by classifying processes as CPU- or I/O-bound based on recent behavior. This brings XINU closer in line with how scheduling is performed in commodity operating systems such as Linux and Windows.

2. Readings

1. [XINU set-up](#)
2. Read Chapters 5-6 of the XINU textbook.

Please use a fresh copy of XINU, `xinu-fall2023.tar.gz`, but for preserving the `helloworld()` function from lab1 and removing all code related to `xsh` from `main()` (i.e., you are starting with an empty `main()` before adding code to perform testing). As noted before, `main()` serves as an app for your own testing purposes. The GTAs when evaluating your code will use their own `main()` to evaluate your XINU kernel modifications.

3. Monitoring process CPU usage and response time [80 pts]**3.1 CPU usage**

We will modify XINU so that it can monitor CPU usage of processes. For a process that has not terminated, its CPU usage will be the time (in unit of msec) that it has spent executing on Galileo's x86 CPU, i.e., in state `PR_CURR`. Estimating CPU usage accurately requires understanding device management (of which hardware clocks are a part) which we will cover later in the course. We will implement a method that is adequate for our purpose of gauging whether our dynamic priority scheduler implemented in Problem 4 achieves fair allocation of CPU cycles to processes.

Introduce a new process table field, `uint32 prtotcpu`, where XINU will record CPU usage (in unit of msec). For a process that is not in state `PR_CURR`, `prtotcpu` will contain XINU's estimate of its total CPU usage from the time of its creation until now, i.e., the moment `prtotcpu` is inspected. For the current process, CPU usage will be the sum of `prtotcpu` and `cputime` where `cputime` is a global variable of type `uint32` to be declared in `system/initialize.c` that estimates the time (in msec) that the current process has spent in state `PR_CURR` after being context-switched in. Modify XINU so that when a context-switch occurs the old process's `prtotcpu` field is updated as `prtotcpu + cputime`. Note that "old process" refers to the process being context-switched out. If `cputime` equals 0, set it to 1 before calculating `prtotcpu + cputime`. For processes that consumed less than 1 msec of user CPU time before context-switching out, we round up CPU usage to 1 msec resulting in overestimation. Reset `cputime` to 0 so that the time that the new process spends in state `PR_CURR` can be monitored.

Since XINU's system timer is set to interrupt every 1 msec, increment `cputime` in `clkdisp.S` similar to how `clkcountermsec` was updated in Problem 4.1 of lab1. Since XINU disables external interrupts when executing upper and lower half kernel code, `cputime` underestimates total CPU usage when XINU executes kernel code. This corresponds to user CPU time in Linux that tracks CPU usage of a process in user mode.

Introduce a new system call, `syscall totcpu(pid32)`, in `system/totcpu.c` that returns total CPU usage in unit of msec of the process specified in the argument. If the argument specifies the PID of the current process, `totcpu()` returns `prtotcpu + cputime`. `prtotcpu` of the current process is updated only when it is context-switched out. If the argument specifies a process that is not in state `PR_CURR`, return `prtotcpu`. Check the PID argument of `totcpu()` to verify that it is valid, and return `SYSERR` if not. Test your implementation to gauge correctness. Describe in `lab3.pdf` your method for doing so. Given that `cputime` is incremented by assembly code in `clkdisp.S`, is it necessary to use the "volatile" qualifier when declaring the global variable in `initialize.c`? Discuss your reasoning.

3.2 Response time

An important metric, especially for I/O-bound processes, is response time (in unit of msec) defined as the time a process resided in XINU's readylist before it became current. Introduce two new process table fields, `uint32 prtoresp`, and `uint32 prtoready`, where `prtoready` initialized to 0 when a process is created counts how many times a process transitioned to state `PR_READY`. When updating `prtoready` only consider XINU process states that we have covered thus far in the course. For example, `PR_RECV` need not be considered since haven't discussed sending and receiving messages. The field `prtoresp` initialized to 0 upon process creation sums the total time a process has spent since its creation waiting in the readylist. We will estimate average response time of a process over its lifetime by dividing `prtoresp` by `prtoready`. Introduce a new system call, `syscall avgresponse(pid32)`, in `system/avgresponse.c` that returns average response time rounded up to an integer (in unit of msec) by dividing `prtoresp` by `prtoready`.

When a process enters XINU's readylist, record the current time since bootloading a backend machine in a new process table field, `uint32 prreadystart`. Reuse the global variable, `uint32 clkcountermsec`, from Problem 4.1, lab1, to keep track of system time elapsed since bootload. When a process transitions from `PR_READY` to `PR_CURR`, subtract `prreadystart` from `clkcountermsec`. If the difference is 0, set the value to 1 (msec), which overestimates response time.

Test your implementation to assess correctness. Describe in `lab3.pdf` your method for doing so. In the case where a newly created process becomes ready for the first time, discuss if the above method for calculating response time incurs inaccuracies. There is no need to rectify any inaccuracies.

Note: When implementing and testing code to monitor CPU usage and response time, use the legacy fixed priority XINU kernel, not the kernel with dynamic priority scheduling in Problem 4. Only after verifying that CPU usage and response time monitoring works correctly under fixed priority scheduling utilize it in Problem 4 for evaluating dynamic priority scheduling.

3.1 CPU Usage

- CPU usage = Time spent executing (i.e. $pr_state == PR_CURR$) while process is still not terminated
- 1. New process table field `uint32 prtotcpu` to record CPU usage in msec ✓
 - ↳ If not PR_CURR , $prtotcpu$ = Moment $prtotcpu$ is inspected - time of creation
 - ↳ Else, $prtotcpu = prtotcpu + cputime$
- 2. Global variable `uint32 cputime` in `system/initialize.c` to estimate how many msec the current process has spent in PR_CURR after being context-switched in ✓
- 3. When context-switch happens, old process' $prtotcpu = prtotcpu + cputime$ ✓
 - ↳ If $cputime == 0$, set it to 1 ✓
- 4. Reset `cputime` to 0 for the new process ✓
 - Only updated when process is context-switched out
- 5. Increment `cputime` in `clkdisp.S` similar to lab1 ✓
- 6. Write syscall `totcpu(pid32)` in `system/totcpu.c` that returns total CPU usage (msec)
 - ↳ If PID is in PR_CURR , return $prtotcpu + cputime$ ✓
 - ↳ Else, return $prtotcpu$ ✓
 - ↳ If PID is not valid, return `YSERR` ✓

3.2 Response Time

- Response time: The time a process resides in XINU's readylist before it becomes current
- 1. Process table field `uint32 prtotready` counts # of times it changes to PR_READY
 - ↳ Initialized to 0 when created ✓
 - ↳ Update for process states we have covered (exclude PR_RCV) → ? → ready.c, resched.c
- 2. Process table field `uint32 prtotresp` sums total time waiting in readylist
 - ↳ Initialized to 0 when created ✓
 - How to update?
- 3. syscall `avgresponse(pid32)` in `system/avgresponse.c`
 - ↳ Returns $prtotresp / prtotready$ rounded up to an integer ✓
- 4. Process table field `uint32 prreadystart` records current time since bootloading a backend machine when the process enters XINU's readylist
- 5. Reuse `uint32 clkcountermsec` from lab1 to keep track of time since bootload ✓
 - ↳ When process changes $PR_READY \rightarrow PR_CURR$, $clkcountermsec - prreadystart$
 - ↳ If difference is 0, set it to 1 ↓ resched.c store in prtotresp?

4. Dynamic priority scheduling [180 pts]

We will use a dynamic process scheduling framework based on process behavior to adaptively modify the priority and time slice of processes as a function of their observed behavior.

4.1 Process classification: CPU-bound vs. I/O-bound

Classification of processes based on observation of recent run-time behavior must be done efficiently to keep the scheduler's overhead to a minimum. A simple strategy is to classify a process based on its most recent scheduling related behavior: (a) if a **process depleted its time slice** the process is viewed as **CPU-bound**; (b) if a **process hasn't depleted its time slice and voluntarily gives up the CPU by making blocking call** we will consider it **I/O-bound**. A third case (c) is a process that is preempted by a higher or equal priority process that was blocked but has become ready. When a process of equal priority becomes ready we preempt the current process. We will remain neutral and not change the priority of the preempted process.

We will repurpose the process table field, `prprio`, to indicate both a priority value and the classification of a process as CPU- or I/O-bound. The former preserves backward compatibility. In the case of (a), `prprio` is set to 1 which classifies a process as CPU-bound. In the case (b), `prprio` is set to 2 which denotes a process as I/O-bound. In the case of (c), the previous value of `prprio` is maintained. We will set `prprio` to 0 to indicate that the process is XINU's idle/null process. The idle process is handled as a special case so that it only runs when there are no current or ready processes in the system. Setting `prprio` to 0 which is strictly less than CPU- and I/O-bound processes achieves that. Define three macros, `#define QUANTUMIO 5`, `#define QUANTUMCPU 50`, `#define QUANTUMIDLE 100`, in `include/process.h` that specify the time slice assigned to I/O-bound, CPU-bound processes, and the idle process, respectively.

4.2 Upper half blocking and lower half preemption

A process is deemed I/O-bound when it **makes a blocking system call in the upper half of the kernel**. The current process may be preempted by another process of equal or higher priority by the actions of the kernel's lower half which responds to interrupts. To keep coding to a minimum, we will use `sleepms()` as a representative blocking system call for all other blocking system calls. Hence in benchmark testing a process is considered I/O-bound if it makes frequent `sleepms()` system calls, thereby voluntarily relinquishing Galileo's CPU before its quantum expires. If a process calls `sleepms()` with an argument that is negative, `sleepms()` will return `YSERR`. An argument value of 0 is allowed in XINU which prompts the kernel to call `resched()`. We will treat this case as I/O-bound behavior since an I/O-bound ready process would preempt the current process.

For preemption events we will only consider those triggered by system timer management in XINU's lower half, `clkhandler()`. The first scenario we will consider is checking if a process in state `PR_SLEEP` needs to be woken up and inserted into the readylist. The second scenario is time slice depletion where `clkhandler()` calls `resched()` to let the scheduler determine who to run next. Note the classification in 4.1 must be implemented by modifying relevant upper and lower half kernel code.

Modify `create()` so that it ignores the third argument and sets the priority of newly created processes to 1. In benchmark testing we will use the process executing `main()` to generate the workload processes and then terminate. Assign to this process priority value 3 by calling `chprio()` from `main()`. Since the process executing `main()` terminates after creating workload processes described in 4.3 it will not affect overall performance evaluation. The global variable `preempt` used to track the current process's remaining time slice must be set to the appropriate value for CPU-bound and I/O-bound processes, as well as for the idle process, in different parts of kernel code.

4.3 Testing and performance evaluation

Perform basic debugging by checking the internal operation of the modified kernel to verify correct operation. Utilize macro `XINUDEBUG` for this purpose which is then disabled to suppress output of debug messages in your submitted code. Use macro `XINUTEST` to enable output of the values during benchmarking where equitable sharing of CPU cycles gauged by CPU usage and improved response time of I/O-bound processes is evaluated.

Benchmark apps *CPU-bound app*. Code a function, `void cpuproc(void)`, in `system/cpuproc.c`, that implements a while-loop. The while-loop checks if `clkcountermsec` exceeds a threshold which is defined by macro `STOPCOND` set to 10000 (msec) in `system/initialize.c`. A process executing `cpuproc()` will terminate when `clkcountermsec` has reached about 10 seconds. By design, a process executing `cpuproc()` hogs Galileo's CPU and is therefore an extreme case of a CPU-bound app.

I/O-bound app. Code a function, `void procio(void)`, in `system/procio.c`, that implements a while-loop to check if `clkcountermsec` has exceeded `STOPCOND`. If so, `procio()` terminates. Unlike the body of `cpuproc()`'s while-loop which is empty, `procio()`'s while-loop body has a for-loop followed by a call to `sleepms()` with argument 50 (msec). Try different values for the bound of the inner for-loop such that it consumes several milliseconds of CPU time. It should not exceed 3 msec but otherwise is not important. The inner for-loop can contain arithmetic operations (even a nested for-loop) to help consume CPU time not exceeding 3 msec. Inspect the value of `clkcountermsec` before and after the for-loop to calibrate the bound.

Benchmark output. Before terminating, `cpuproc()` and `procio()` print PID, "CPU-bound" or "I/O-bound" depending on type, CPU usage by calling `totcpu()`, response time by calling `avgresponse()`, and `clkcountermsec`. Modify `nulluser()` so that in the infinite while-loop it checks if `clkcountermsec` has exceeded `STOPCOND`. If so, `nulluser()` outputs PID, "idle process", CPU usage, response time, and `clkcountermsec`. Code `nulluser()` so that it does this only once. As before `nulluser()` does not terminate.

Workload scenarios We will consider homogenous and mixed workloads in benchmarks A-C. Benchmark D considers mixed workloads where starvation may occur.

Benchmark A. Spawn a workload generation process using `create()` that runs `main()` which calls `chprio()` to set its priority to 3. The process running `main()` spawns 4 app processes each executing `cpuproc()`. Call `resume()` to ready the 4 processes after creating them. The workload generation process terminates after creating and resuming the four benchmark processes. Output by the four app processes upon termination at around 10 seconds of wall time should indicate approximately equal sharing of CPU time and similar response times. Some CPU time will have been consumed by the idle process which will also have a response time. The total sum of CPU times should be approximately 10 seconds. Discuss your results in `lab3.pdf`.

Benchmark B. Repeat benchmark scenario A with the difference that the app processes execute `procio()`. Since the apps are homogenous, their CPU usage and response time should be similar. Discuss your finding in `lab3.pdf`. Compare the results of benchmarks A and B.

Benchmark C. Let the workload generator `main()` create 8 app processes, half of them executing `cpuproc()`, the other half `procio()`. The four CPU-bound processes should get similar CPU usage and response times since they are homogenous. The same goes for the four I/O-bound processes. Across the two groups, CPU-bound processes should get both higher CPU usage and response times than I/O-bound processes. Discuss your finding in `lab3.pdf`.

Benchmark D. Code an app, `void parasite(void)`, in `system/parasite.c` that **exploits a weakness of the dynamic priority XINU scheduler which classifies a process that calls `sleepms(0)` as exhibiting I/O-bound behavior**. The logic underlying `parasite()` is to occupy the CPU for slightly less than `QUANTUMIO`

4.1 Process Classification

- CPU-bound: Process depleted its time slice
- I/O-bound: Hasn't depleted but voluntarily gives up the CPU by making a blocking call
- 3rd case: Process preempted by a higher/equal priority process that was blocked but is ready
- `prlib prprio` is 1 for CPU-bound, 2 for I/O-bound, 0 for XINU's idle/null process
- `#define QUANTUMCPU 50`
 `QUANTUMIO 5` in `include/process.h` ✓
 `QUANTUMIDLE 100`

4.2 Upper half blocking & lower half preemption

- Use `sleepms()` as a representative blocking system call
 - ↳ Process is considered I/O-bound if it calls `sleepms()` frequently
 - ↳ Calling `sleepms()` with negative argument returns `YSERR` ✓
 - ↳ Argument 0 prompts kernel to call `resched()` — called by `yield()`
- Only consider preemption events triggered by `clkhandler()`
 - ↳ Check if a process in `PR-SLEEP` needs to be woken up & inserted into readylist
 - ↳ Time slice depletion where `clkhandler()` calls `resched()`
- Modify `create()` so it ignores `arg3` & set the priority of new processes to 1 ✓
- For benchmark testing, use `main()` to generate workload processes & then terminate
 - ↳ Assign `main()`'s priority to 3 by calling `chprio()` ✓
- Global variable `preempt` used to track current process's remaining time slice must be set appropriately for CPU/I/O-bound processes & idle processes in the kernel code

4.3 Testing & performance evaluation

- Use `XINUDEBUG` to suppress output of debug messages in submitted code.
Use `XINUTEST` to enable output of the values during benchmarking
- Code `void cpuproc(void)` in `system/cpuproc.c` that implements an empty while loop ✓
 - ↳ While loop checks if `clkcountermsec` exceeds `STOPCOND` (set to 1000 msec in `include/process.h`)
 - ↳ A process executing this terminates after 10 seconds ✓
- Code `void procio(void)` in `system/procio.c`
 - ↳ Has a while loop to check if `clkcountermsec` exceeds `STOPCOND`. If so, terminate ✓
 - ↳ Body of while loop has for-loop followed by `sleepms()` with argument 50 ✓
 - ↳ Try different bounds such that it consumes several ms (not more than 3) ✓
- Both functions print PID, 'CPU'/'I/O' + '-bound', CPU usage (`totcpu()`), response time (`avgresponse()`) & `clkcountermsec` ✓

- Modify `nulluser()` so in infinite while-loop, it checks if `clkcounter msec` exceeds `STOPCOND`
 - ↳ If so, output PID, "idle process", CPU usage, response time & `clkcounter msec` ✓
 - ↳ Code it so it only does it once ✓
- Benchmark A:
 - ↳ Process running `main()` spawns 4 app processes each executing `cpuproc()` ✓
 - ↳ Call `resume()` to ready the 4 processes after creating them ✓
 - ↳ Process running `main()` terminates after creating & resuming the 4 processes
 - ↳ Output upon termination should indicate approx. equal sharing of CPU time & similar response times.
 - ↳ Total sum of CPU time should be approx 10 seconds ✓
- Benchmark B:
 - ↳ Repeat A but with `procio()`
 - ↳ Should have similar CPU usage & response time
- Benchmark C:
 - ↳ `main()` create 8 processes, half `cpuproc()` half `procio()` ✓
 - ↳ Similar process should have similar CPU usage & response times ✓
 - ↳ CPU-bound processes should get higher CPU usage & response times ✓
- Benchmark D:
 - ↳ Code `void parasite(void)` in `system/parasite.c`
 - ↳ To occupy the CPU for slightly less than `QUANTUMIO` & call `sleepms(0)`
 - ↳ Appears to kernel as being I/O-bound when it has CPU-bound behavior
 - ↳ Modify `yield()` such that current process gets fresh time slice `QUANTUMIO` & priority is set to 2 ✓
 - ↳ Create 3 processes executing `cpuproc()` & one executing `parasite()`
 - ↳ `parasite()` should have most of CPU time

Bonus

- Code `procio()`

and call `sleepms(0)` before the kernel has a chance to potentially preempt the process. If all processes in the readylist are CPU-bound, `parasite()` will get a fresh time slice of `QUANTUMIO` and continue to run. Hence `parasite()` appears to the kernel as being I/O-bound when its actual behavior is CPU-bound. Create three processes that execute `cpuproc()` and one process that executes `parasite()`. After all processes terminate at around 10 seconds the process executing `parasite()` should have received the lion's share of CPU time. Discuss your results in `lab3.pdf`.

Bonus problem [20 pts]

Modify the I/O-bound app, `procio()`, into a variant, `prociogang()`, in `system/prociogang.c` so that a group of I/O-bound processes can easily starve a CPU-bound process. Create three processes executing `prociogang()` and one process executing `procpu()`. When they terminate at around 10 seconds of wall clock time, the three I/O-bound processes should have received approximately equal performance while starving the CPU-bound process. Describe the changes you made in `prociogang()` and discuss your results.

Note: The bonus problem provides an opportunity to earn extra credits that count toward the lab component of the course. It is purely optional.

Turn-in instructions

General instructions:

When implementing code in the labs, please maintain separate versions/copies of code so that mistakes such as unintentional overwriting or deletion of code is prevented. This is in addition to the efficiency that such organization provides. You may use any number of version control systems such as GIT and RCS. Please make sure that your code is protected from public access. For example, when using GIT, use `git` that manages code locally instead of its on-line counterpart `github`. If you prefer not to use version control tools, you may just use manual copy to keep track of different versions required for development and testing. More vigilance and discipline may be required when doing so.

The TAs, when evaluating your code, will use their own test code (principally `main()`) to drive your XINU code. The code you put inside `main()` is for your own testing and will, in general, not be considered during evaluation.

If you are unsure what you need to submit in what format, consult the [TA notes](#) link. If it doesn't answer your question, ask during PSOs and office hours which are scheduled M-F.

Specific instructions:

1. Format for submitting written lab answers and `kprintf()` added for testing and debugging purposes in kernel code:

- Provide your answers to the questions below in `lab3.pdf` and place the file in `lab3/`. You may use any document editing software but your final output must be exported and submitted as a pdf file.
- For problems where you are asked to print values using `kprintf()`, use conditional compilation (C preprocessor directives `#define` combined with `#ifdef` and `#endif`) with macro `XINUTEST` (in `include/process.h`) to effect print/no print depending on if `XINUTEST` is defined or not. For your debug statements, do the same with macro `XINUDEBUG`.

2. Before submitting your work, make sure to double-check the [TA Notes](#) to ensure that any additional requirements and instructions have been followed.

3. Electronic turn-in instructions:

i) Go to the `xinu-fall2023/compile` directory and run "make clean".

ii) Go to the directory where `lab3` (containing `xinu-fall2023/` and `lab3.pdf`) is a subdirectory.

For example, if `/homes/alice/cs354/lab3/xinu-fall2023` is your directory structure, go to `/homes/alice/cs354`

iii) Type the following command

```
turnin -c cs354 -p lab3 lab3
```

You can check/list the submitted files using

```
turnin -c cs354 -p lab3 -v
```

Please make sure to disable all debugging output before submitting your code.

[Back to the CS 354 web page](#)