

CS 354 Fall 2023

Lab 5: Asynchronous Handling of Events using Callback Function [270 pts]

Due: 11/16/2023 (Thu.), 11:59 PM

1. Objectives

The goal is to provide kernel support for asynchronous event handling using callback functions that preserve isolation/protection when events are triggered by asynchronous events.

2. Readings

- Read Chapters 10, 13-14 of the XINU textbook.
-

Please use a fresh copy of XINU, `xinu-fall2023.tar.gz`, but for preserving the `helloworld()` function from lab1 and removing all code related to `xsh` from `main()` (i.e., you are starting with an empty `main()` before adding code to perform testing). As noted before, `main()` serves as an app for your own testing purposes. The GTAs when evaluating your code will use their own `main()` to evaluate your XINU kernel modifications.

3. Asynchronous event handling using callback function [270 pts]

3.1 Overview

In Problem 5, lab4, we introduced kernel support for run-time detour to a user callback function specified in the argument of `dreamvac()` when `sleepms()` returned. Instead of the callback function being called directly by user code, ROP was used by the kernel to modify the return address of `sleepms()` to arrange the detour. In this problem we will support asynchronous user callback function execution in response to asynchronous events.

3.2 Callback function registration

Code a system call, `syscall cbsignal(uint16 etype, void (* cbf) (void), uint32 val)`, in `system/cbsignal.c` that can be used by processes to register a user callback function, `cbf`, with the kernel. If registration is successful, the kernel will arrange execution of the callback function when the event specified by the first argument `etype` occurs. To preserve isolation/protection, the kernel will arrange execution of `cbf` in user mode and in the context of the process that registered the callback function. In XINU, user mode is interpreted to mean when a process returns from a system call or from an interrupt where there are clear dividing lines between kernel code and user code. The third argument, `val`, is meaningful only for some event types. `cbsignal()` returns 0 if the specified arguments are valid and the request can be met. It returns `SYSERR` otherwise. Validity is determined per event type type described below.

3.3 Alarm event

Registration. If `etype` is 1 then the kernel is asked to set an alarm event, a timer specified by the third argument `val` (in unit of msec). We will reuse the millisecond counter, `clkcountermsec`, from lab1 to support alarm events.

Introduce four new process table fields, `uint16 pretype`, `void (* prcbf1)()`, `uint32 pralarmtime`, `uint16 pralarmreg`, where `pretype` and `pralarmreg` are initialized to 0. When `cbsignal(1, myfunc, val)` is called, the kernel returns `SYSErr` if `val` is greater than 5000. That is, we will not support timers longer than 5 seconds. `cbsignal()` returns `SYSErr` if `pralarmreg` is not 0, meaning that we allow a process to have one outstanding timer alarm only. `cbsignal()` returns `SYSErr` if the function pointer, `cbf`, falls outside the text segment memory region of XINU. Check `initialize.c` which utilizes the boundary addresses of XINU's text segment.

If the arguments of `cbsignal()` are verified to be valid, `cbsignal()` sets `pralarmreg` to 1 indicating that an alarm event has been registered for the current process. `prcbf1` is set to `myfunc`. `pralarmtime` is set to `clkcountermsc + val`.

ROP detour. To reduce coding we will consider one canonical case -- lower half system timer -- where the kernel checks if the current process has an alarm event that has expired for which a detour needs to be arranged. Each time that `clkhandler()` is invoked in response to XINU's 1 msec clock interrupt, `clkhandler()` will check if `clkcountermsc` is greater than or equal to `pralarmtime` of the current process. If so, a detour needs to be arranged. Instead of jumping to the callback function `prcbf1` directly, the kernel will set up the process's stack so executing `iret` in `clkdisp()` will cause a jump to system function, `void cbfmanager(void)`, to be coded in `system/cbfmanager.S`. `cbfmanager()` is a kernel function that does not need to be executed in kernel mode, hence we will refer to it as a system function. `cbfmanager()` calls the user callback function, `prcbf1`, which then returns to `cbfmanager()`. When `cbfmanager()` executes `ret` it jumps to the original return address of `clkdisp()`.

Recall from lab2 that when an interrupt occurs x86 pushes values contained in `EFLAGS`, `CS`, `EIP` onto the current process's stack. To induce `cbfmanager()` to return to the original return address of `clkdisp()`, use a different method from Problem 5, lab4, where `clkdisp()` rearranges its stack -- upon checking that a detour to `cbfmanager()` needs to be made -- so that it jumps to `cbfmanager()` upon executing `iret`. Note that when `clkhandler()` returns to `clkdisp()` the current process's stack contains the 8 registers saved by `pushal` followed by the values of `EIP`, `CS`, `EFLAGS` saved by x86 when a clock interrupt occurred. The first step is to shuffle the values near the top of the stack so that the 8 register values are located below `EFLAGS`, `CS`, `EIP` at the top of the stack. Then savekeep the original return address `EIP` before overwriting it with function pointer `cbfmanager`. What it means to "savekeep" is up to you to determine. When `iret` is executed, the values of `cbfmanager`, `EFLAGS`, `CS` are popped by x86 upon jumping to `cbfmanager()`. When `cbfmanager()` starts executing the top of the stack must contain the 8 register values.

One of the reasons we are introducing `cbfmanager()` and coding it in assembly is because of the need to restore the register values saved by `clkdisp()` before `cbfmanager()` returns to the original return address. In Problem 5, lab4, this was not a concern since `sleepms()`, its caller, and `vacation()` were all coded in C. That is, per CDECL the caller of `sleepms()` is responsible for saving/restoring `EAX`, `ECX`, `EDX`, and the callee `sleepms()` for saving/restoring the rest. When `sleepms()` jumps to `vacation()` all the register values would be in the same state as when the caller of `sleepms()` had called `sleepms()`. Hence `vacation()`, per CDECL, will save/restore all register values but for `EAX`, `ECX`, `EDX` which the caller of `sleepms()` will restore when `vacation()` jumps to it. For the alarm event in lab5 it is the responsibility of `cbfmanager()` to restore state before returning to the original return address when executing `ret`. Also, be mindful of additional details such as whether the `ESP` value saved by `pushal` in `clkdisp()` is still valid or needs to be changed before `cbfmanager()` executes `popal` followed by `ret`.

Discuss in lab5.pdf your method for modifying `clkdisp.S` and coding `cbfmanager.S` so that the detour is correctly facilitated, including saving/restoring of register values and updating `pralarmreg`. The bottom line is correct execution of the interrupted process when `cbfmanager()` returns (i.e., jumps) to the original return address of `clkdisp()`.

Note: When kernels implement trapped system calls following the traditional approach as in lab2, the system call dispatcher in the lower half needs to perform ROP based detour arrangement similar to `clkdisp()`. Also, if an alarm expires when a process has been context-switched out, it is the scheduler who decides when the process becomes current. Until such time, execution of the callback function is delayed.

3.4 Asynchronous message receive event

Registration. If `etype` is 2 then the kernel is asked to invoke user callback function `cbf` if a message arrives from a sender. In our uniprocessor Galileo x86 backends, for a message to arrive a sender process must be running on the CPU who calls `send()`. Hence the receiver process cannot be current, i.e., has been context-switched out. The third argument of `cbsignal()`, `val`, is ignored. As in the alarm event, function pointer `cbf` is verified to fall within the text segment of XINU. If `etype` is not 1 or 2, `cbsignal()` returns `YSERR`. The process table field, `uint16 pretype`, is set to 2 if a callback function for an alarm event is not outstanding. `pretype` is set to 3 indicating that both an alarm event and an asynchronous message event are outstanding. When an asynchronous message receive event or alarm event has been handled, `pretype` must be updated accordingly.

Introduce new process table fields, `void (* prcbf2)()`, `uint16 prmsgreg`, where `prmsgreg` is initialized to 0. `prcbf2` is used to store the function pointer `cbf`. `prmsgreg` is set to 1 if there is an outstanding asynchronous message receive event. Only one outstanding asynchronous message receive event is allowed. — `YSERR` if `prmsgreg != 0`

ROP detour. To simplify coding, we will consider the receiver calling `sleepms()` as the canonical case where the receiver process has been context-switched out when the sender runs. Modify `send()` so that it checks if the receiver has a callback function registered for asynchronous message receive. If so, use the technique from Problem 5, lab4, to locate the return address of `sleepms()` to its caller in the stack of the receiver process. Modify the receiver's stack so that a detour to the receiver's callback function `prcbf2` is arranged. When `prcbf2` returns it jumps to the original return address of `sleepms()`. Thus ROP-based detour arrangement is a reuse of the implementation of Problem 5, lab4, albeit performing stack modification in `send()` as opposed to `wakeup()` or in `resched()` before context-switching in a process.

Update of `prmsgreg` must await until the receiver process wakes up, becomes ready, and is selected by the scheduler to become current. Only then is `prmsgreg` updated since the receiver is about to make a detour to its callback function when `sleepms()` returns. During testing we will ignore the scenario where the receiver wakes up and becomes current without having received a message during sleep. This can be handled by utilizing a sender flag which is not needed for our purposes.

3.5 Testing

Test and verify that your implementation works correctly. Describe in lab5.pdf your test procedure to gauge correctness. Please make sure to adequately annotate your code which will count for 10% of the total score.

Bonus problem [40 pts]

Extend Problem 3.4 so that the canonical case includes an app blocking by calling `suspend()` on itself, in addition to support for `sleepms()`. Assume that a second process calls `resume()` with the PID of the suspended process to unblock it. Put your code in a separate subdirectory `bonus/` under `xinu-fall2023/`. Test and verify that your code works correctly.

Note: The bonus problem provides an opportunity to earn extra credits that count toward the lab component of the course. It is purely optional.

Turn-in instructions

General instructions:

When implementing code in the labs, please maintain separate versions/copies of code so that mistakes such as unintentional overwriting or deletion of code is prevented. This is in addition to the efficiency that such organization provides. You may use any number of version control systems such as GIT and RCS. Please make sure that your code is protected from public access. For example, when using GIT, use git that manages code locally instead of its on-line counterpart github. If you prefer not to use version control tools, you may just use

manual copy to keep track of different versions required for development and testing. More vigilance and discipline may be required when doing so.

The TAs, when evaluating your code, will use their own test code (principally `main()`) to drive your XINU code. The code you put inside `main()` is for your own testing and will, in general, not be considered during evaluation.

If you are unsure what you need to submit in what format, consult the [TA notes](#) link. If it doesn't answer your question, ask during PSOs and office hours which are scheduled M-F.

Specific instructions:

1. Format for submitting written lab answers and `kprintf()` added for testing and debugging purposes in kernel code:

- Provide your answers to the questions below in `lab5.pdf` and place the file in `lab5/`. You may use any document editing software but your final output must be exported and submitted as a pdf file.
- For problems where you are asked to print values using `kprintf()`, use conditional compilation (C preprocessor directives `#define` combined with `#if` and `#endif`) with macro `XINUTEST` (in `include/process.h`) to effect print/no print depending on if `XINUTEST` is defined or not. For your debug statements, do the same with macro `XINUDEBUG`.

2. Before submitting your work, make sure to double-check the [TA Notes](#) to ensure that any additional requirements and instructions have been followed.

3. Electronic turn-in instructions:

i) Go to the `xinu-fall2023/compile` directory and run `"make clean"`.

ii) Go to the directory where `lab5` (containing `xinu-fall2023/` and `lab5.pdf`) is a subdirectory.

For example, if `/homes/alice/cs354/lab5/xinu-fall2023` is your directory structure, go to `/homes/alice/cs354`

iii) Type the following command

```
turnin -c cs354 -p lab5 lab5
```

You can check/list the submitted files using

```
turnin -c cs354 -p lab5 -v
```

Please make sure to disable all debugging output before submitting your code.

[Back to the CS 354 web page](#)

